

Evaluation — L1-inception-requirements-extractor-evaluator

Quality Gates

- ☐ All findings have category + description + fix_applied (or unfixable_reason)
- ☐ Score per dimension is 0.0–1.0
- ☐ Scores are post-fix (reflect corrected document state)
- ☐ Verdict is pass (all $\geq$ threshold) or fail
- ☐ If fixes applied, corrected PRD re-uploaded
- ☐ No hallucinated fixes (only corrects genuine issues)
- ☐ Faithfulness check uses original input as ground truth

Scores

Dimension	$\geq$	Checks
Accuracy	0.90	Findings are genuine issues, not false positives
Fix quality	0.85	Fixes resolve issue without introducing new problems
Completeness	0.90	All rubric dimensions evaluated, all inline gates checked

Reflection Checklist

- ☐ Every finding traces to a specific rubric criterion from KB or inline gates
- ☐ Fixes don't invent requirements not in original input
- ☐ Pass/fail verdict matches post-fix scores (not contradictory)
- ☐ Re-uploaded artifact is the corrected version, not original
- ☐ Unfixable items have clear reason why fix is impossible