

Evaluation — L1-inception-vision-generator-evaluator

Quality Gates

- ☐ All findings have category + description + fix_applied
- ☐ Score per dimension is 0.0–1.0
- ☐ Overall verdict is pass (all scores $\geq$ threshold) or fail
- ☐ If fixes applied, corrected artifact re-uploaded
- ☐ No hallucinated fixes (only corrects what's genuinely wrong)

Scores

Evaluator	$\geq$	Checks
Accuracy	0.90	Findings are genuine issues, not false positives
Fix quality	0.85	Fixes resolve the issue without introducing new problems
Completeness	0.90	All rubric dimensions evaluated

Reflection Checklist

- ☐ Every finding traces to a specific rubric criterion
- ☐ Fixes don't invent content not grounded in domain KB
- ☐ Pass/fail verdict matches scores (not contradictory)
- ☐ Re-uploaded artifact is the corrected version, not original