

Evaluation Criteria — L1-inception-workflow-summariser

Quality Gates (must pass)

Criterion	Threshold	Method
All agents mentioned	100% of input agents appear in execution_flow	Automated: count check
No invented data	0 claims not traceable to input	LLM-as-Judge: faithfulness
Plain English	No raw JSON in output	Automated: regex check
Status accuracy	Workflow outcome matches agent statuses	Automated: logic check
Resources consolidated	No duplicate KB/guardrail/tool entries	Automated: uniqueness

Evaluation Scores (LLM-as-Judge)

Evaluator	Threshold	Direction
Faithfulness	≥ 0.95	Every claim traces to an agent's execution_summary
Hallucination	≤ 0.05	No invented counts, artifacts, or statuses
Completeness	≥ 0.95	All agents, artifacts, and resources covered
Clarity	≥ 0.85	Non-technical reader can understand
Consistency	≥ 0.90	No contradictions between sections

Quality Rubric

Dimension	Score 9-10	Score 7-8	Score 5-6	Score < 5
Coverage	All agents, artifacts, resources captured	1 minor item missed	Several items missed	Major gaps
Clarity	Plain English, concise, stakeholder-ready	Mostly clear, minor jargon	Technical language mixed in	Unreadable by non-technical
Accuracy	Every fact verified against input	1 minor inaccuracy	Several inaccuracies	Invented data
Structure	All 5 sections present and logical	Sections present but slightly disorganised	Missing sections	Unstructured

Reflection Checklist

The agent must self-verify before delivering:

- ☐ Every agent in the input array is mentioned in `execution_flow`
- ☐ `request_and_intent` reflects the actual workflow trigger (not invented)
- ☐ Artifact names/locations match what agents actually produced
- ☐ Resources (KBs, guardrails, tools) are consolidated with no duplicates
- ☐ Failed agents are highlighted in the outcome
- ☐ No raw JSON or schema fragments in the output
- ☐ Summary is concise (10-20 bullet points, not verbose paragraphs)

Reflection Process (mandatory)

1. **Generate** initial summary from agent outputs
2. **Log** `[REFLECTING]` Checking output against evaluation criteria
3. **Check** every item in the Reflection Checklist above
4. **Identify** missing agents, invented data, or unclear language
5. **Log** each finding: `[REFLECTING]` Found: `<description>`
6. **Fix** each issue — amend the output silently
7. **Log** each resolution: `[REFLECTING]` Resolved: `<what was fixed>`
8. **Deliver** only the final corrected output